

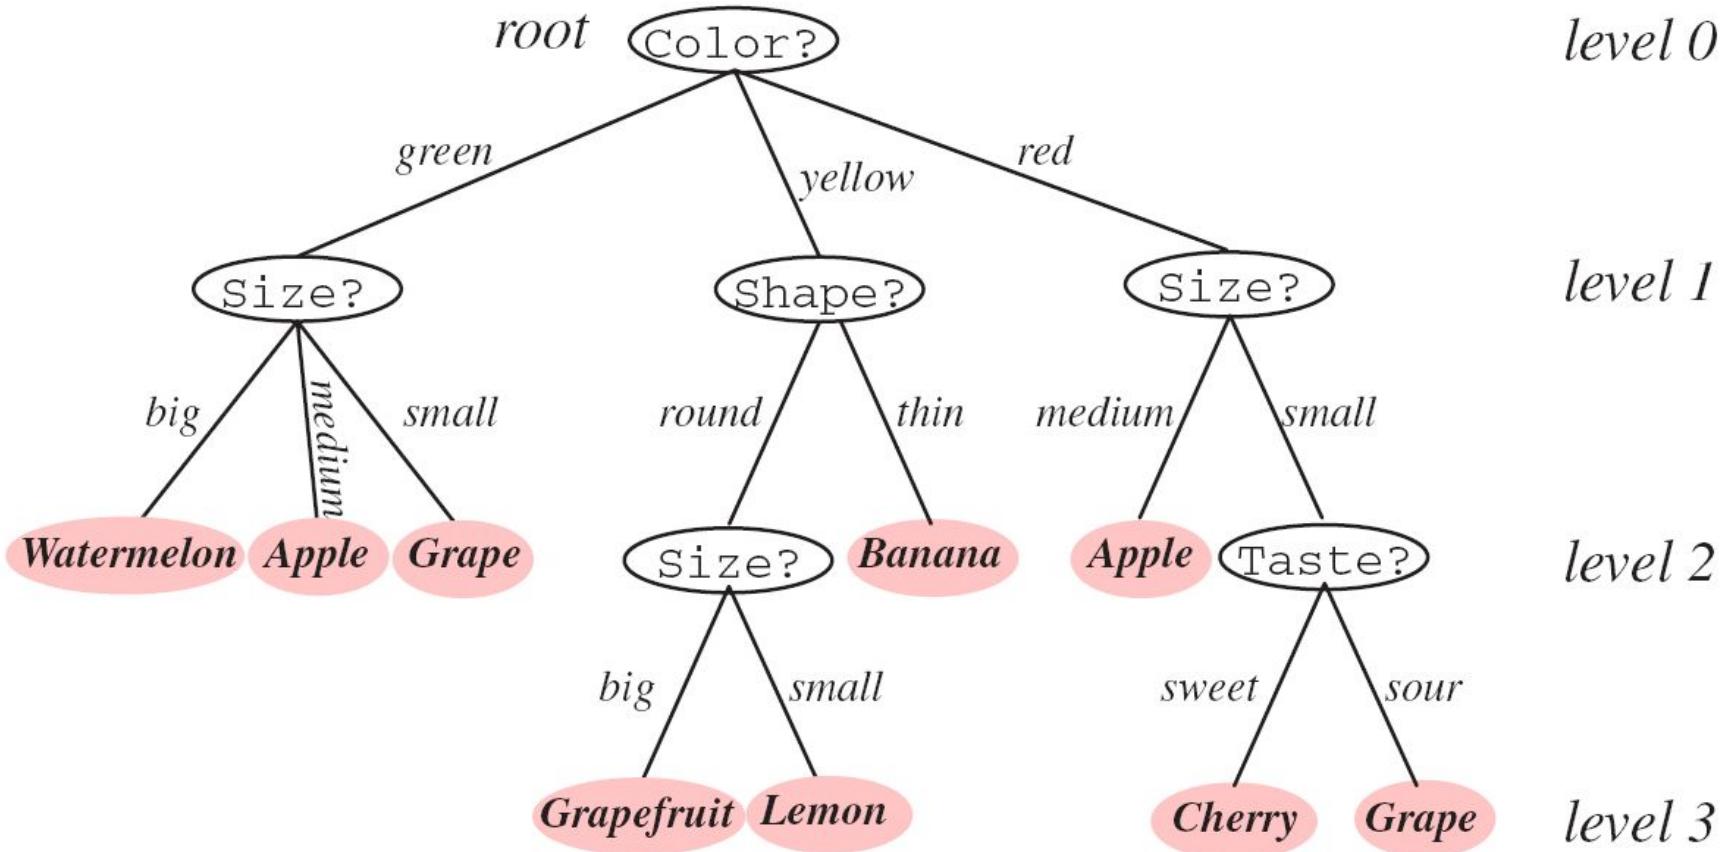
Decision Trees and Random Forests

Assoc. Prof. Sascha Hornauer

Center for Robotics
Mines Paris
PSL Université

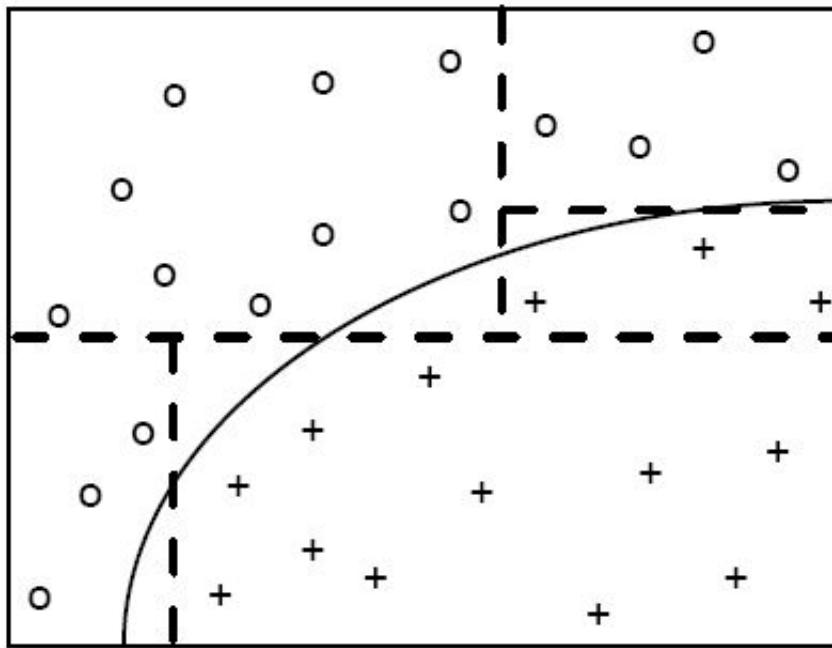
sascha.hornauer@minesparis.psl.eu

What is a Decision Tree?



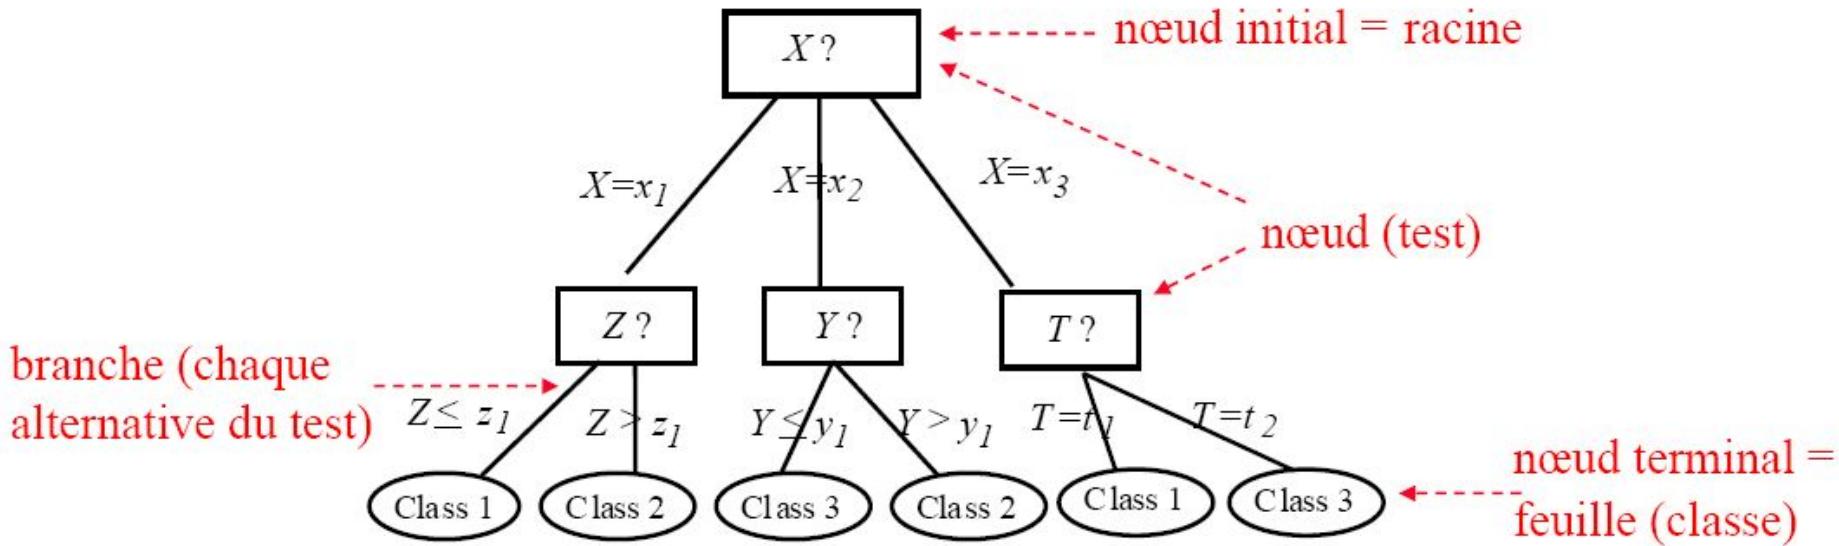
Classification by a tree of tests

General principle of Decision Trees



Classification by sequences of tests organized in a tree, and corresponding to a *partition of input space into class-homogeneous sub-regions*

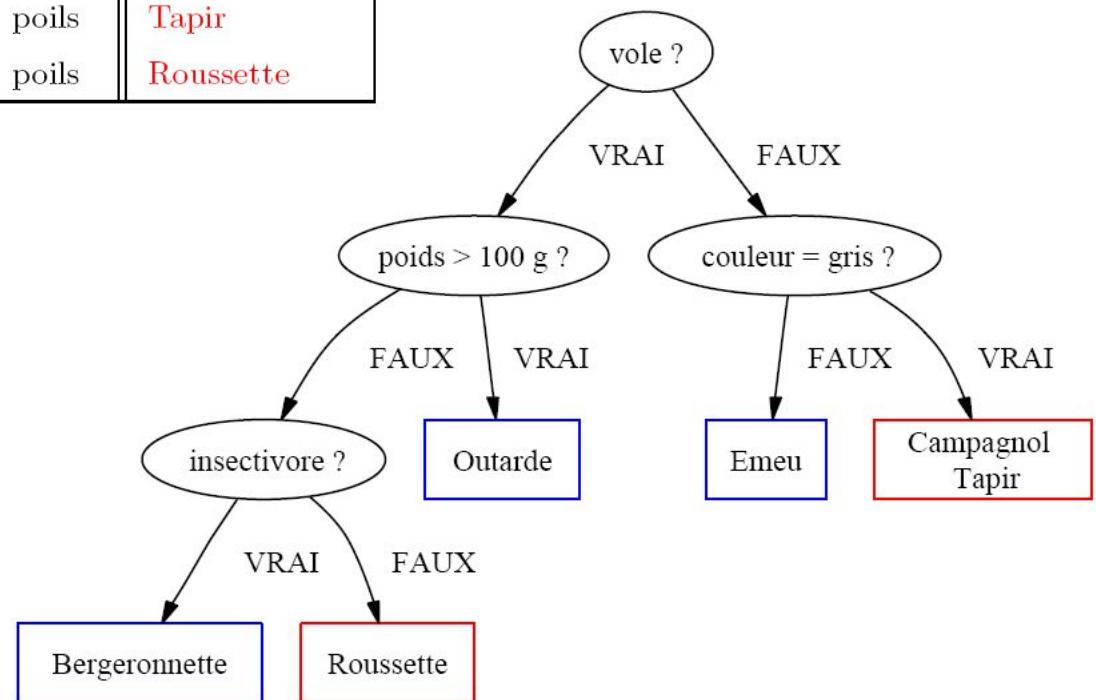
Example of Decision Tree



- **Classification rule:** go from root to a leaf by evaluating the tests in nodes
- **Class of a leaf:** class of the majority of training examples “arriving” to that leaf

“Induction” of the tree?

vole	poids	couleur	alimentation	peau	animal
OUI	1 kg	roux	granivore	plumes	Outarde
OUI	20 g	gris et jaune	insectivore	plumes	Bergeronnette
NON	100 kg	noir et blanc	omnivore	plumes	Emeu
NON	5 g	gris	granivore	poils	Campagnol
NON	40 kg	gris	herbivore	poils	Tapir
OUI	60 g	noir	frugivore	poils	Roussette



Is it the best tree??

- Exhaustive search in the set of all possible trees is computationally intractable

□ Recursive approach to build the tree:

build-tree(X)

IF all examples "entering" X are of same class,
THEN build a leaf (labelled with this class)
ELSE

- choose (using some criterion!) the BEST (attribute;test) couple to create a new node
- this test splits X into 2 sub-trees X_1 and X_r
- build-tree(X_1)
- build-tree(X_r)

Criterion for choosing attribute and test

- **Measure of heterogeneity of candidate node:**
 - entropy (ID3, C4.5)
 - Gini index (CART)
- **Entropy:** $H = -\sum_k (p(w_k) \log_2(p(w_k)))$ with $p(w_k)$ probability of class w_k (estimated by proportion N_k/N)
 - minimum ($=0$) if only one class is present
 - maximum ($=\log_2(\#\text{of}\text{classes})$) if equi-partition
- **Gini index:** $\text{Gini} = 1 - \sum_k p^2(w_k)$

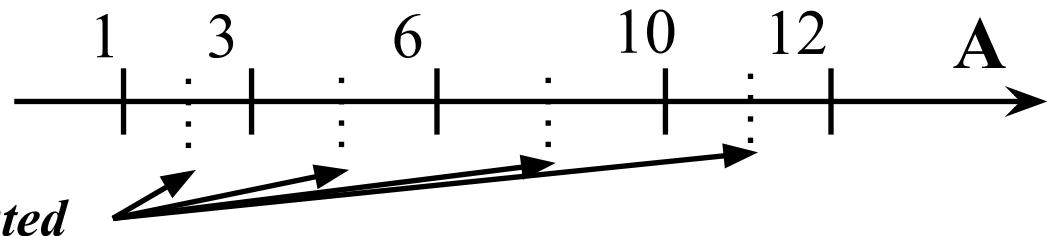
Homogeneity gain by a test

- Given a test T with m alternatives and therefore orienting from node N into m “sub-nodes” N_j
 - Let $I(N_j)$ be the heterogeneity measures (entropy, Gini, ...) of sub-nodes, and $p(N_j)$ the proportions of elements directed from N towards N_j by test T
-
- the homogeneity gain brought by test T
is $\text{Gain}(N, T) = I(N) - \sum_j p(N_j) I(N_j)$
 - Simple algo = choose the test maximizing this gain
(or, in the case of C4.5, the “relative” gain $G(N, T)/I(N)$, to avoid bias towards large m)

Tests on *continuous-valued attributes*

- Training set is **FINITE** idem for the # of values taken **ON TRAINING EXAMPLES** by any attribute, even if continuous-valued
- In practice, examples are sorted by increasing value of the attribute, and only $N-1$ potential threshold values need to be compared (typically, the medians between successive increasing values)

For example, if values of attribute A for training examples are 1;3;6;10;12, the following potential tests shall be considered:
 $A > 1.5; A > 4.5; A > 8; A > 11$)



- “Obvious” stopping rules:
 - all examples arriving in a node are of same class
 - all examples arriving in a node have equal values for each attribute
 - node heterogeneity stops decreasing
- Natural stopping rules:
 - # of examples arriving in a node < minimum threshold
 - Control of generalization performance (on independent validation set)
- A posteriori pruning: remove branches that are impeding generalization (bottom-up removal from leaf while generalization error does not decrease)

Criterion for a posteriori pruning of the tree

Let T be the tree, v one of its nodes, and:

- $IC(T, v) = \# \text{ of examples Incorrectly Classified by } v \text{ in } T$
- $IC_{ela}(T, v) = \# \text{ of examples Incorrectly Classified by } v \text{ in } T' = T \text{ pruned by changing } v \text{ into a leaf}$
- $n(T) = \text{total \# of leaves in } T$
- $nt(T, v) = \# \text{ of leaves in the sub-tree below node } v$

THEN the criterion chosen to minimize is:

$$w(T, v) = (IC_{ela}(T, v) - IC(T, v)) / (n(T) * (nt(T, v) - 1))$$

□ Take simultaneously into account
error rate and tree complexity

Pruning algorithm

Prune(T_{\max}) :

$K \leftarrow 0$

$T_k \leftarrow T_{\max}$

WHILE T_k has more than 1 node, DO

FOR EACH node v of T_k DO

compute $w(T_k, v)$ on train. (or valid.) examples

END FOR

choose node v_m that has minimum $w(T_k, v)$

$T_{k+1} : T_k$ where v_m was replaced by a leaf

$k \leftarrow k + 1$

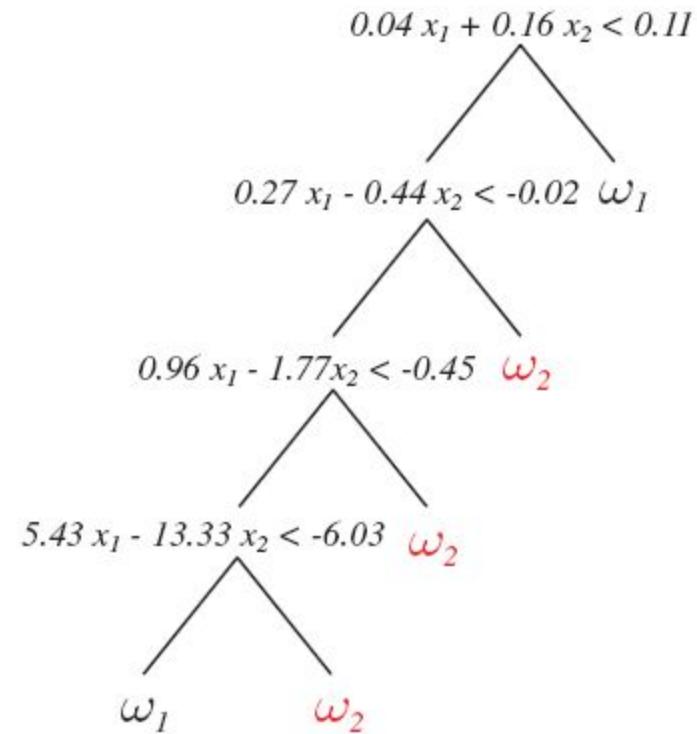
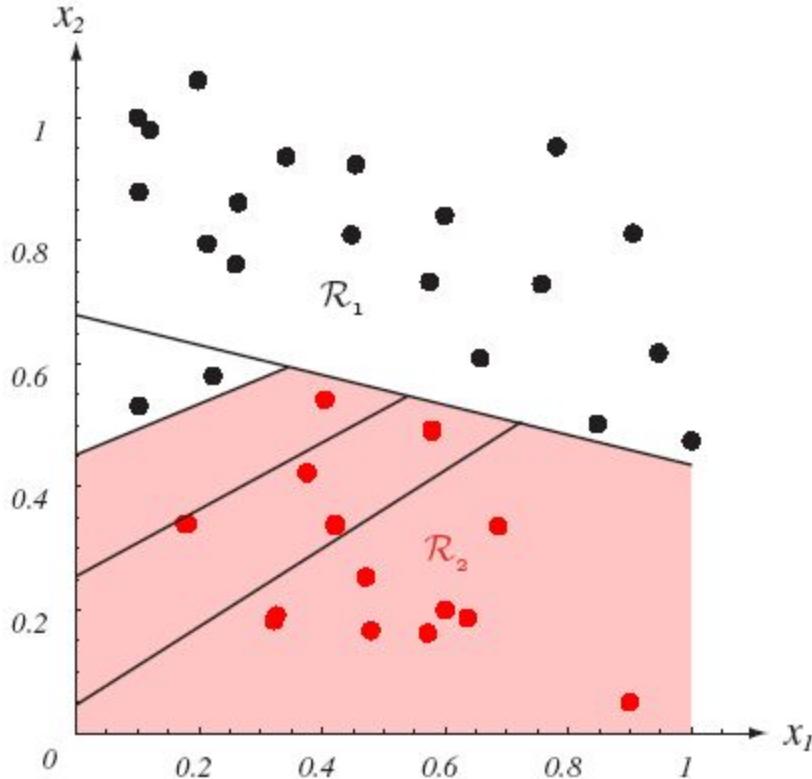
END WHILE

Finally, select among $\{T_{\max}, T_1, \dots, T_n\}$ the pruned tree that has the smallest classification error on the validation set

Names of variants of Decision Tree variants

- **ID3 (Inductive Decision Tree, Quinlan 1979):**
 - only “discrimination” trees (i.e. for data with all attributes being qualitative variables)
 - heterogeneity criterion = entropy
- **C4.5 (Quinlan 1993):**
 - Improvement of ID3, allowing “regression” trees (ie continuous-valued attribute), and handling missing values
- **CART (Classification And Regression Tree, Breiman et al. 1984):**
 - heterogeneity criterion = Gini

Other variant: multi-variate trees



- Homogeneity criterion (entropy or Gini)
- Recursion stop criteria:
 - Maximum depth of tree
 - Minimum # of examples associated to each leaf
- Pruning parameters

Pros and cons of Decision Trees

- **Advantages**
 - Easily manipulate “symbolic”/discrete-valued data
 - OK even with variables of varying amplitudes
(no need for explicit normalization)
 - Multi-class **BY NATURE**
 - **INTERPRETABILITY** of the tree!
 - Identification of “important” inputs
 - Very efficient classification (especially for very-high dimension inputs)
- **Drawbacks**
 - High sensitivity to noise and “erroneous outliers”
 - Pruning strategy rather delicate

Principle: “*Strength lies in numbers*”
[en français, “*L’union fait la force*”]

- A forest = a set of trees
- **Random Forest:**
 - Train a large number T (~ few 10s or 100s) of *simple* Decision Trees
 - Use a *vote of the trees* (majority class, or even estimates of class probabilities by % of votes) if classification, or an *average of the trees* if regression

Algorithm proposed in 2001 by Breiman & Cutler

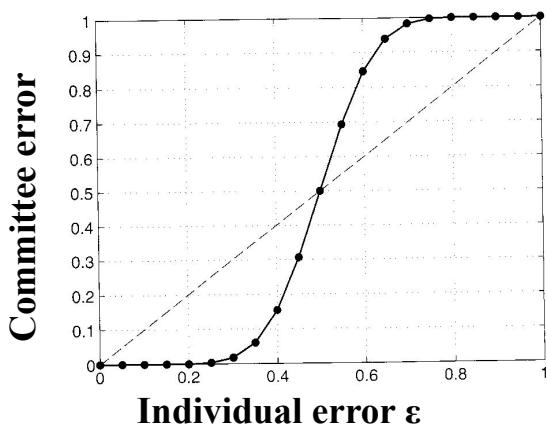
Theoretical background of “ensembling” methods

Set-up a “committee of experts”
 each one can be wrong, but combining opinions
 increases the chance to obtain correct prediction!

Theoretical justification:

- suppose N independent classifiers, each with same error rate $E_{\text{gen}} = \varepsilon$
- decision by a “majority” vote is wrong if and only if more than half of the committee is wrong

$$\square \quad \text{Error}_{\text{committee}} = \sum_{k=N/2}^N C_k^N \varepsilon^k (1 - \varepsilon)^{N-k}$$



**Spectacular improvement of decision
 (under condition that $\varepsilon < 0.5$!!)...
 ...and the larger N (# of experts),
 the bigger the improvement**

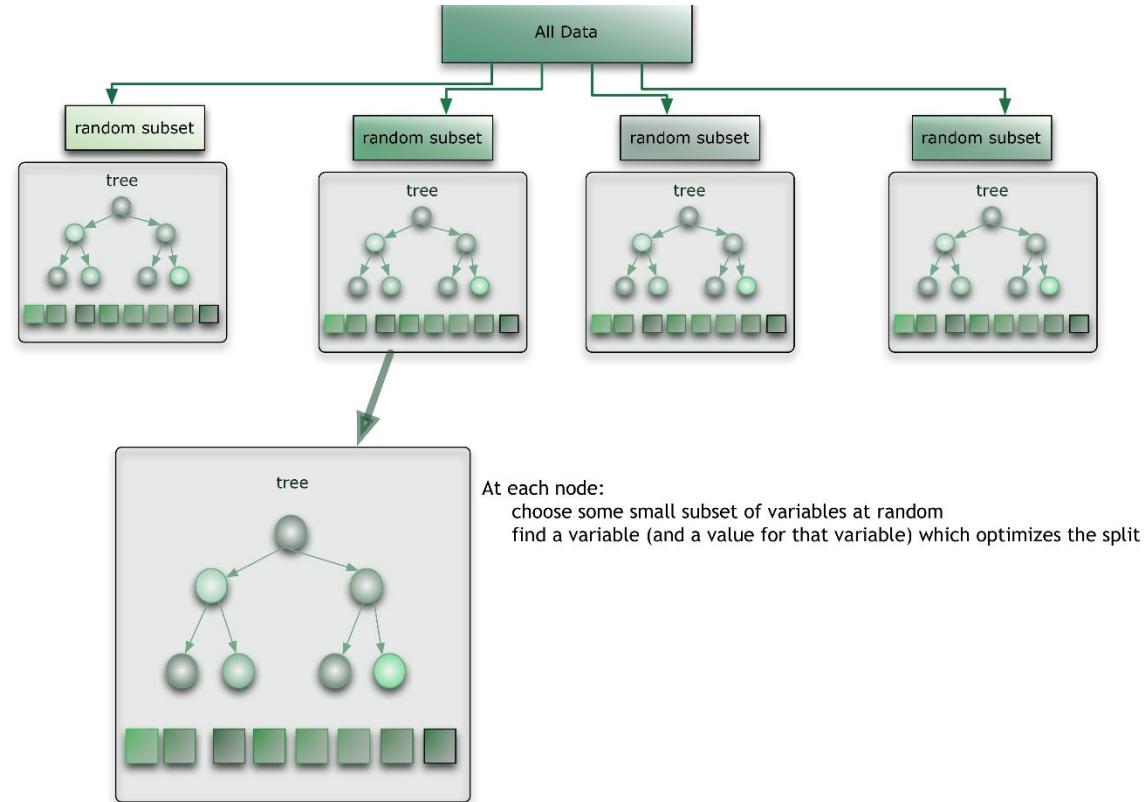
“wisdom of the crowd” (?)

Learning of a Random Forest

Goal= obtain trees as decorrelated as possible

- each tree is learnt on a random different subset (~2/3) of the whole training set

- each node of each tree is chosen as an optimal “split” among only k variables randomly chosen from all d inputs (and $k \ll d$)



- Each tree is learnt using CART *without pruning*
- The maximum depth p of the trees is usually strongly limited ($\sim 2 \text{ à } 5$)

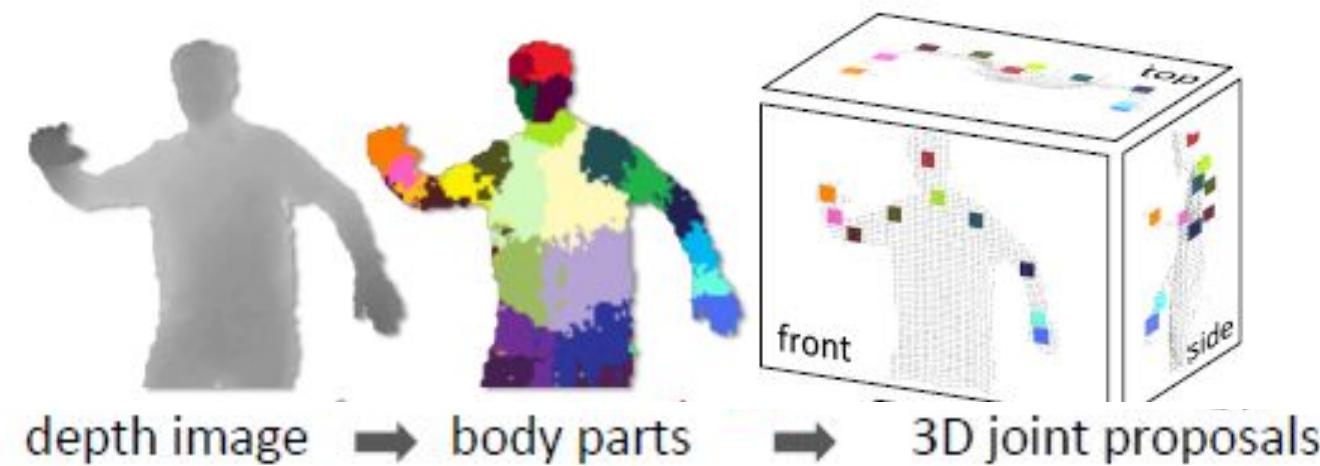
$Z = \{(x_1, y_1), \dots, (x_n, y_n)\}$ training set,
each x_i of dimension d

FOR $t = 1, \dots, T$ ($T = \#$ of trees in the forest)

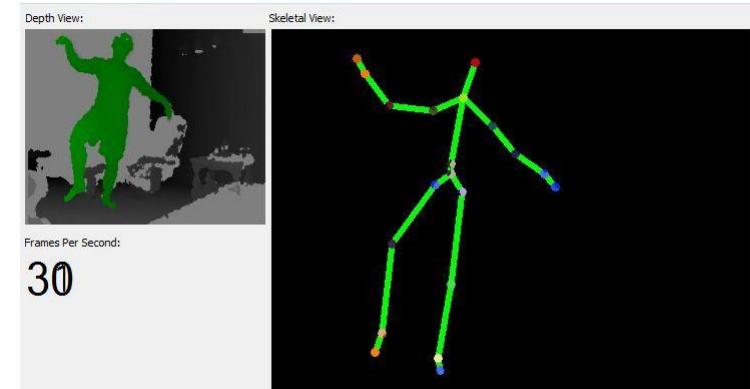
- Randomly choose m examples in Z ($\square Z_t$)
- Learn a tree on Z_t , with CART modified for randomizing variables choice: each node is searched as a test on one of ONLY k variables randomly chosen among all d input dimensions ($k \ll d$, typically $k \sim \sqrt{d}$)

RdF "Success story"

“Skeletonization” of persons (and movement tracking) with Microsoft Kinect™ depth camera



*Algo of Shotton et al.
using RdF for
labelling body parts*



Hyper-parameters for Random Forests

- The number of trees
- Maximum depth of trees
- The size of randomized subset of training examples
- The proportion K/D of attributes considered for inference of each tree

Pros and Cons of Random Forests

- **Advantages**
 - **VERY FAST** recognition
 - **Multi-class by nature**
 - **Efficient on large-dimension inputs**
 - **Robustness to outliers**
- **Drawbacks**
 - Training often rather long
 - Extreme values often incorrectly estimated in case of regression