

# **Boosting:** **combining elementary classifiers** **to learn a “strong” classifier**

Assoc. Prof. Sascha Hornauer#

Center for Robotics

Mines Paris

PSL Université

`Fabien.Moutarde@minesparis.psl.eu`

`http://people.minesparis.psl.eu/fabien.moutarde`

# Essential principle: “wisdom of the crowd”

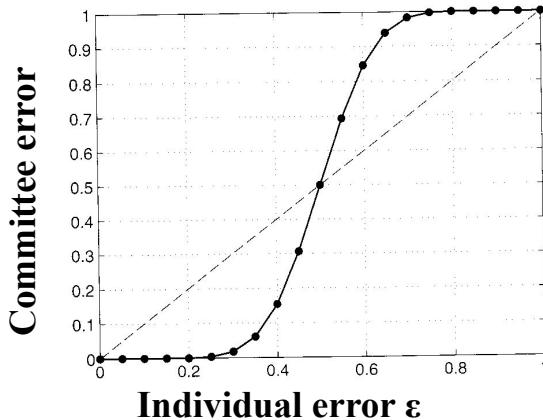
## Set-up a “committee of experts”

each one can be wrong, but combining opinions increases the chance to obtain correct prediction!

## Theoretical justification:

- suppose  $N$  independent classifiers, each with same error rate  $E_{\text{gen}} = \varepsilon$
- decision by a “majority” vote is wrong if and only if more than half of the committee is wrong

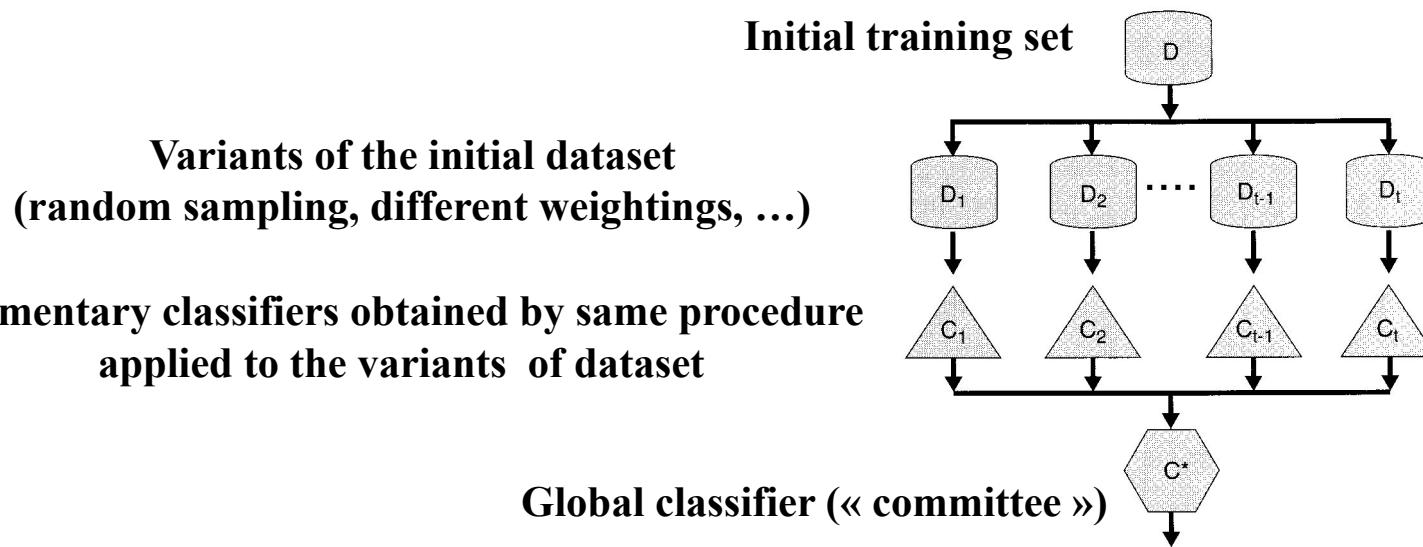
□  $\text{Error}_{\text{committee}} = \sum_{k=N/2}^N C_k^N \varepsilon^k (1 - \varepsilon)^{N-k}$



Spectacular improvement of decision  
(under condition that  $\varepsilon < 0.5!!$ )...  
...and the larger  $N$  (# of experts),  
the bigger the improvement

# Various methods to produce the elementary classifiers to combine

- Use totally different algorithms
- Same algorithm, but with different parameters and/or initializations
- **Modify the training set**



**Very GENERAL methods,**  
**applicable to enhance any « elementary » algorithm**

# Bagging

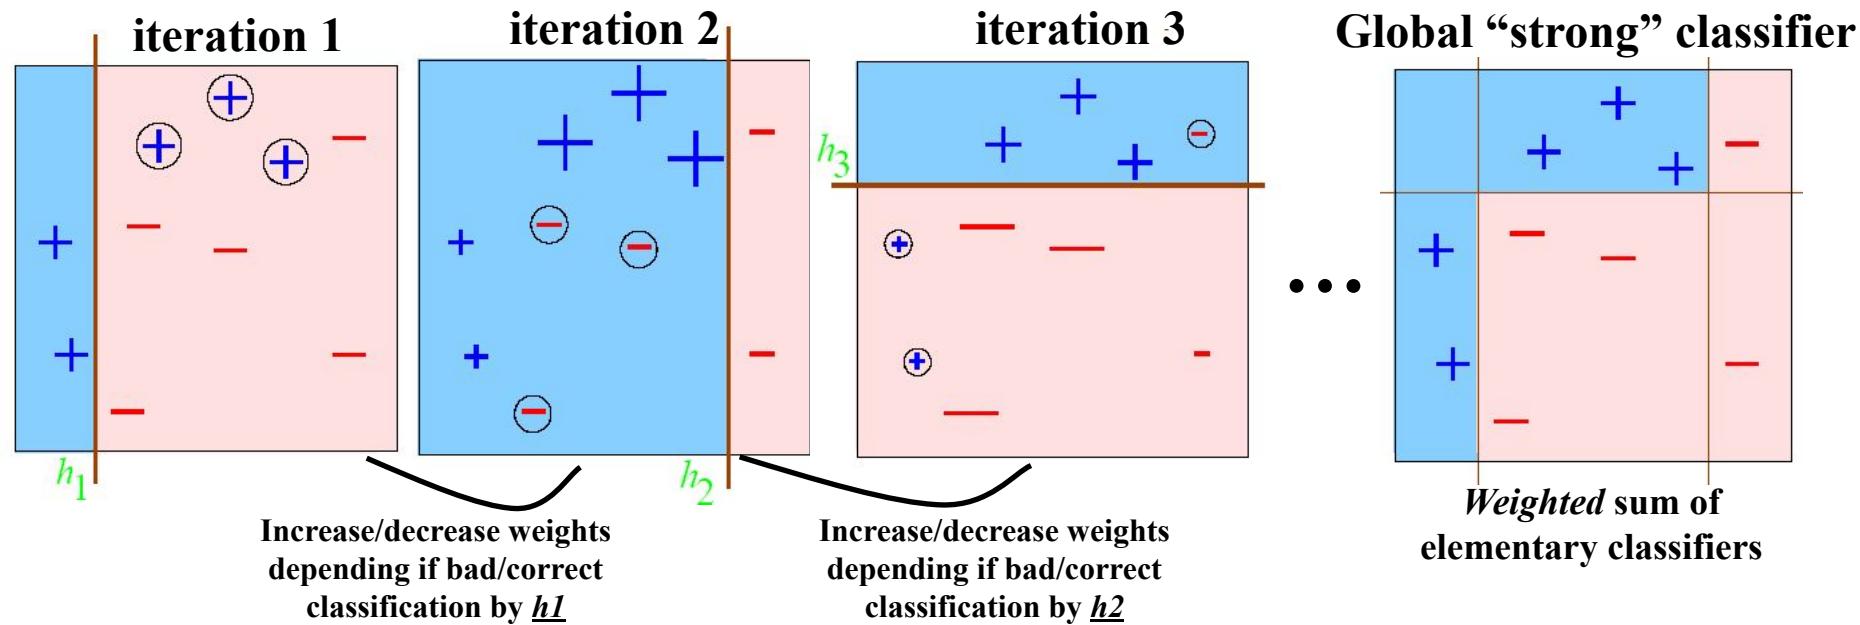
**Variants of training set obtained by  
random sampling (with re-placement)  
from initial dataset**

(kind of "***bootstrap***"  random duplication/erasure of  
some examples, depending on the variant)

- Useful and efficient in particular if the “elementary” algorithm is “sensitive to data noise” (because then different variants of training set shall induce quite different classifiers)
- Reduces *over-fitting*, because the final classifier is a kind of average of classifiers learnt on different realizations of the same data

# Boosting

Iterative method for adding new classifiers to the committee:  
**variants of training dataset obtained by successive weightings of the same examples**  
 (computed for “focusing” on hard examples,  
 i.e. incorrectly classified by previous elementary classifiers)



## adaBoost (“adaptive Boosting”)

- Initial training set:

$S = \{(x_1, u_1), \dots, (x_k, u_k)\}$ , with  $u_i \in \{+1, -1\}$ ,  $i=1, k$

- Initial weights:  $w_0(x_i) = 1/m$  for all  $i=1, k$  (or  $1/2p$  for pos, and  $1/2n$  for neg)

- For each iteration (or round)  $t$  from 1 to  $T$ , do:

- Learn/choose 1 classification rule  $h_t$  on  $(S, w_t)$  using algorithm A

- Compute weighted error  $\varepsilon_t$  of  $h_t$  on  $(S, w_t)$ :  $\varepsilon_t = \sum_{i=1}^k w_t(x_i) \times \|h_t(x_i) - u_i\|$

- Deduce reliability score  $\alpha_t$  of  $h_t$ :  $\alpha_t = \frac{1}{2} \ln \left( \frac{1 - \varepsilon_t}{\varepsilon_t} \right)$  [ $\alpha_t > 0$  if  $\varepsilon_t < 0.5$ , and  $\rightarrow +\infty$  if  $\varepsilon_t \rightarrow 0$ ]

- Modify weights of examples, i.e. for  $i$  from 1 to  $k$ , do:

$$w_{t+1}(x_i) = \frac{w_t(x_i)}{Z_t} \times \begin{cases} e^{-\alpha_t} \text{ si } h_t(x_i) = u_i \text{ (i.e. } x_i \text{ bien classé)} \\ e^{+\alpha_t} \text{ si } h_t(x_i) \neq u_i \text{ (i.e. } x_i \text{ mal classé)} \end{cases}$$

- Output the global “strong” classifier:  $H(x) = \operatorname{signe} \left( \sum_{t=1}^T \alpha_t h_t(x) \right)$

Freund & Schapire (inventors of the algorithm)  
have demonstrated the following theorem:

**If each elementary classifier has error-rate <0.5,  
then empirical error of  $H_T$  on S decreases  
exponentially with the number T of iterations**

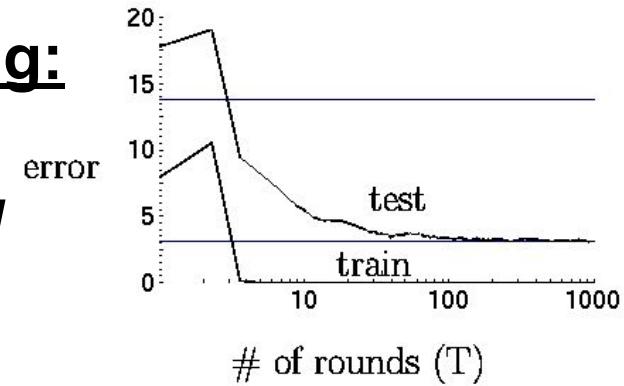
**More precisely**  $E_{emp}(H_T) = \frac{1}{k} \sum_{i=1}^k \|H_T(x_i) - u_i\|$  **is bounded by:**

$$E_{emp}(H_T) \leq \prod_{t=1}^T \left[ 2\sqrt{\varepsilon_t(1-\varepsilon_t)} \right] = \prod_{t=1}^T \sqrt{1 - 4\gamma_t^2}$$

**(where  $\gamma_t = 0.5 - \varepsilon_t$  is the improvement of  $h_t$  compared to random decision)**

# Boosting and margins

Typical error training curve for boosting:  
**the generalization error continues to decrease many iterations after training error becomes zero!!**

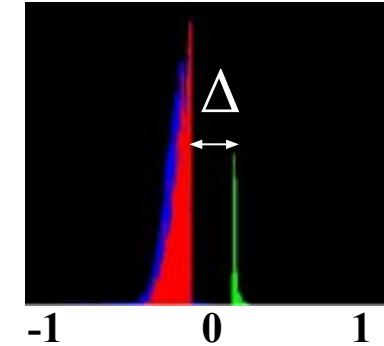


**Reason: even after training\_error reaches 0,  
adaBoost continues to increase *margins*  
i.e. output ≠ between negative and positive examples**

**Margin  $m$  of strong classifier  $H_T$  on example  $x_i$ :**

$$m(H_T, x_i) = u_i \sum_{t=1}^T \alpha_t h_t(x_i) / \sum_{t=1}^T \alpha_t$$

$m(x_i) \in [-1;+1]$ , and  $x_i$  correctly classified  $\Leftrightarrow m(x_i) > 0$ ,  
but **the more  $|m|$  increases, the larger the  $\Delta$  separation  
between positive and negative examples**



# Risk of over-fitting by adaBoost?

- Weight increase of « ambiguous » examples
  - risk of over-fitting?
- Fortunately, generalization error bounded by:

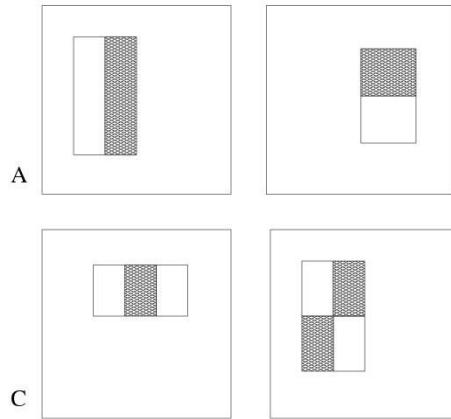
$$E_{gen}(H_T) < \Pr(m(H_T, x) \leq \theta) + O\left(\sqrt{\frac{\delta}{n\theta^2}}\right)$$

where n is the # of examples, and  $\delta$  the VC-dimension of  $h_t$  family

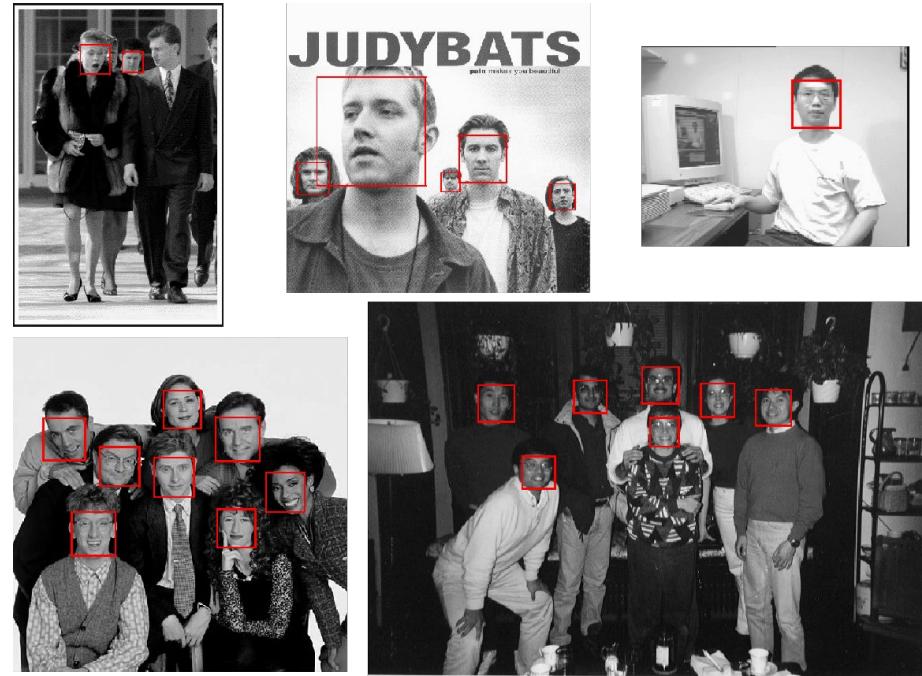
- if  $p(m(H_T, x) < \theta)$  very low for a big-enough  $\theta$ ,  
then good generalization.

In practice the margin m increases with iterations,  
so this bound decreases ☺

- Visual object detection by selection-boosting of « Haar features »; initial example initial = face detection by Viola&Jones (2001)



**Weak classifiers =  
comparison of sums of  
pixels in adjacent rectangles**



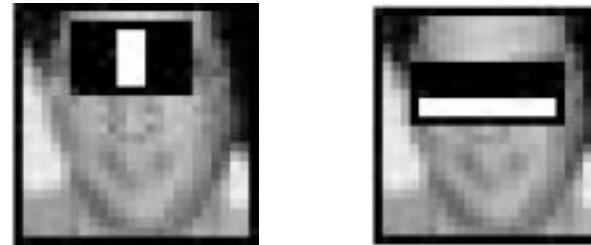
**Result of applying strong classifier on  
multiple sub-windows of various sizes  
and positions (“window scanning”)**

# Boosting as feature selection (and weighting)

**adaBoost = weighted vote by a committee of "weak classifiers" obtained by iterative weightings of examples**

Final **STRONG** classifier:  $H(x) = \text{sign} \left( \sum_{t=1}^T \alpha_t h_t(x) \right)$

Idea of Viola&Jones in 2001: use as weak classifier very simple boolean features selected in a family (e.g. all Haar-like features)  
 ⇔ Weak Learner = search of feature with lowest weighted error



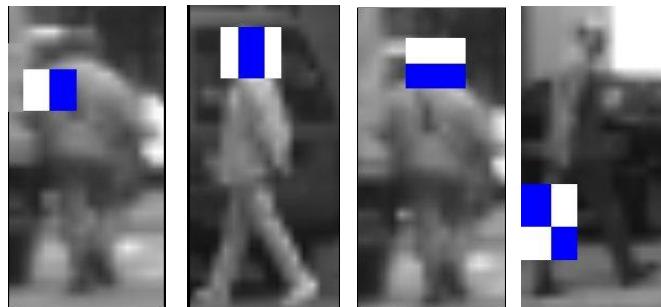
*Using a 24x24 pixels detection window, with all possible combinations of horizontal&vertical location and scale of Haar, the full set of features has 45,396 ≠ features (and ~10 times more in a 32x32 window)  brute-force exhaustive search possible!*

# Outcome of boosting with $\neq$ feature families



Typical connected-Control-Points selected during Adaboost training

For comparison, typical Adaboost-selected Haar features



# Result of car & pedestrian detection with boosting



**Cars (from behind) : ~ 95% detection  
with < 1 false alarm / image**



**Pedestrian (daytime) : ~80% detection  
with < 2 false alarms / image**

[Research conducted in ~2009  
@ center for Robotics  
of MINES ParisTech]

# Hyper-parameters for adaBoost

- Type of weak classifiers assembled
- The “weak learner”  $L$  which trains/generates a new weak classifier at each iteration (and potential hyper-parameters of  $L$ )
- # of iterations (= also the # of assembled weak classifiers)

- **Advantages**
  - Can boost the performance of ANY learning algo (if able to handle weighting of examples)
  - Can build a strong classifier with ANY type of very weak classifiers (slightly better than random)
  - Can be used as an algorithm for selecting “weakly discriminative features” (cf. Viola & Jones)
- **Drawbacks**
  - Training time can be rather long (especially in “discriminative-feature selection” case)
  - Potential risk of over-fitting?